

Risk Analysis of Pension Portfolio Liabilities in the Netherlands

Kerem Başol

April 18, 2026

Abstract

This report provides valuation of the pension plan liabilities for ISSurance, focusing on the Green, Silver and Gold portfolios within the Netherlands. The analysis begins by exploring the historical evolution of mortality across both time and age. It is followed by the calibration of Lee-Carter model to analyze the a_x , b_x , and k_t parameters. To take lifespans exceeding standard data limits into account, the Kannisto method is applied to extrapolate mortality rates and determine the period life expectancy and Expected Present Value (EPV) of the liabilities. The scope is then expanded from period to cohort thinking. This reveals critical differences in valuation when future mortality improvements are explicitly modeled. Later, a holistic risk assessment is achieved through the simulation of mortality paths and the application of Solvency II stress tests for mortality and financial shocks. These simulations allow for a deeper assessment of the uncertainty surrounding the best estimate of total liabilities. Finally, the report concludes by projecting the total portfolio cashflows over time. This establishes the liquidity required to cover liabilities under various scenarios including stress scenarios.

Contents

1	Historical Evolution of Mortality in the Netherlands	3
1.1	Data and its Structure	3
1.2	Calculation of Central Death Rates	3
1.3	Visualizing Historical Trends	4
1.4	Fitting the Lee-Carter Model	4
1.5	Old-Age Mortality: The Kannisto Method	6
1.6	Evaluation of Period Life Expectancy	7
1.7	Financial Assumptions: EIOPA Risk-Free Rate	8
1.8	Valuation of Period Liabilities	9
2	Period to Cohort Thinking	10
2.1	Cohort Life Expectancy	10
2.2	Valuation of Cohort Liabilities	12
2.3	Longevity and Mortality Shocks and their Effects	13
2.4	Financial Shocks: Interest Rate Sensitivity	15
2.5	Stochastic Interest Rates: Vasicek Model	17
3	Stochastic Mortality Scenarios	18
3.1	Simulating Future Mortality	18
3.2	Life Expectancy Distributions	19
3.3	Stochastic Liability Valuation	20
3.4	Age-Dependent Capital Adequacy	22
4	Portfolio Valuation	23
4.1	Individual Cashflow Projections	23
4.2	Aggregate Portfolio Cashflows	23
4.3	Total Liability Valuation and Uncertainty	24

4.3.1	Capital Requirements and Risk Margin	25
4.4	Sensitivity Analysis: Biometric vs. Financial Shocks	25
4.4.1	Longevity and Mortality Shock	25
4.4.2	Interest Rate Sensitivity	26
5	Conclusion and Recommendations	27
A	Proof of Peak Expected Present Value at Retirement Age	28
B	Source Code and Reproducibility	28

1 Historical Evolution of Mortality in the Netherlands

1.1 Data and its Structure

To build the valuation of pension liabilities, it is essential to first understand the underlying mortality dynamics of the target population. To achieve this, we utilize the Lee-Carter model [1], which provides a robust framework for forecasting future mortality trends and generating realistic future scenarios. The model is specified as:

$$\ln(m_{x,t}) = a_x + b_x k_t + \epsilon_{x,t} \quad (1)$$

In this specification, a_x represents the average mortality level at age x , and k_t tracks the general mortality trend over time. The b_x parameter measures how sensitive mortality at age x is to changes in the overall trend, while $\epsilon_{x,t}$ accounts for residual variation.

Because the central death rate ($m_{x,t}$) is the fundamental input required to calibrate the Lee-Carter model, I have sourced historical mortality data for the Netherlands from the Human Mortality Database (HMD) [2]. Specifically, I utilize 1×1 tables, representing one-year age intervals and one-year time intervals, for both death counts ($D_{x,t}$) and exposure-to-risk ($E_{x,t}$).

Inspection of the raw data reveals records spanning from the year 1850 up to 2023, covering ages from 0 to 110+.

```
> head(nld_deaths)
  Year Age Female Male Total
1 1850  0   8735 10804 19539
2 1850  1   2292  2454  4746
3 1850  2   1009  1055  2064
4 1850  3    591   634  1225
5 1850  4    461   521   982
6 1850  5    411   425   836
...
> tail(nld_deaths)
  Year Age Female Male Total
19309 2023 105  34.30 4.29 38.59
19310 2023 106  18.77 2.18 20.95
19311 2023 107   9.46 1.00 10.46
19312 2023 108   4.72 0.27  4.99
19313 2023 109   2.30 0.00  2.30
19314 2023 110+   1.32 0.00  1.32
```

1.2 Calculation of Central Death Rates

The raw counts of deaths ($D_{x,t}$) and exposures ($E_{x,t}$) must be converted into central death rates, denoted as $m_{x,t}$. The central death rate represents the average risk of death for an individual aged x exactly, during the calendar year t . It is defined as the ratio of total deaths to the total exposure to risk, in this case exposure is the person-years lived.

$$m_{x,t} = \frac{D_{x,t}}{E_{x,t}} \quad (2)$$

As $m_{x,t}$ is not a standard probability, it can take greater values than 1. Because the denominator represents *person-years* rather than the initial number of individuals, the central death rate can exceed 1. For example, consider a group of 10 individuals observed at the start of the year. If all 10 pass away within the first month, they collectively contribute only $10 \times \frac{1}{12} \approx 0.833$ person-years

of exposure. Consequently, the central death rate would be $m_{x,t} = 10/0.833 = 12$, which is greater than 1.

To compute our empirical rates, we join the deaths and exposures datasets. We strictly filter out the "110+" age category at this stage. Removing this open-ended interval ensures the `Age` variable remains numeric.

1.3 Visualizing Historical Trends

To explore the evolution of mortality across different generations in our portfolio, we isolate specific decade ages (30 through 90) and plot the natural logarithm of their central death rates, $\log(m_{x,t})$, over time. The resulting visualization is presented in Figure 1 below.




Figure 1: Historical evolution of log-mortality rates in the Netherlands (1850–2023) across ages.

Commentary on the Evolution of Mortality:

Observing the plotted $\log(m_{x,t})$ curves, several key demographic trends become apparent:

- **General Downward Trend:** Across all age groups, there is a clear, long-term downward trajectory in mortality rates. This reflects continuous historical improvements in healthcare and living standards in the Netherlands.
- **Age Stratification:** The lines are distinctly parallel and stacked chronologically by age. As expected, base mortality is strictly increasing with age; a 90-year-old consistently has a higher baseline risk of death than a 30-year-old, regardless of the calendar year.
- **Historical Shocks:** The plots capture significant biometric shocks, most notably the sharp upward spikes corresponding to the 1918 Spanish Flu pandemic and the periods surrounding the World Wars. A slight uptick in recent years (around 2020–2021) is also visible, reflecting the impact of the COVID-19 pandemic.
- **Varying Rates of Improvement:** The slope of improvement is not identical across all ages. Younger ages (e.g., 30 and 40) show steeper historical declines in mortality compared to older ages (e.g., 80 and 90), indicating that medical advancements have historically had a large relative impact on preventing early-to-mid-life deaths.

1.4 Fitting the Lee-Carter Model

To calibrate the Lee-Carter model, the raw data must be reshaped into a matrix format (years by ages) required by the modeling packages (e.g., `StMoMo`). For the initial calibration phase, we restrict the age range to 0 – 90. Mortality rates for extreme old ages (90+) are often volatile due to low

exposure and will be addressed separately using the Kannisto extrapolation method. The following implementation prepares the death ($D_{x,t}$) and exposure ($E_{x,t}$) matrices and fits the standard Lee-Carter model:

```
> Dxt[1:5, 1:6]
  1850 1851 1852 1853 1854 1855
0 19539 21058 24000 21248 20749 22228
1  4746  5549  5825  5138  5479  6913
2  2064  2723  2557  2582  2820  3863
3  1225  1317  1394  1497  1674  2364
4   982   926   805   986  1060  1663
...
> Ext[87:91, 170:174]
  2019  2020  2021  2022  2023
86 56446.92 57311.46 59116.00 59981.57 61718.86
87 50798.45 50394.30 51000.94 52893.15 53866.87
88 44520.35 44662.15 44052.29 44807.17 46743.62
89 37206.62 38472.43 38378.07 37917.50 38702.46
90 30747.90 31498.30 32395.15 32395.21 31970.28
```

Once the model is fitted, it decomposes the historical mortality surface into three distinct parameters: a_x , b_x , and k_t . To present these findings, we extract the parameters from the fitted object and visualize them in Figures 2, 3, 4.

Comments on the Lee-Carter Parameters:

The calibration yields three specific insights into the demographic profile of the Netherlands:

- **Baseline Profile (a_x):** This parameter represents the average log-mortality across all years. As expected, it demonstrates a "tick" shape: relatively high infant mortality, dropping to a minimum during late childhood and subsequently rising exponentially with advanced age.
- **Improvement Sensitivity (b_x):** This indicates which age groups benefit most when the general mortality trend improves. The higher values observed at younger ages (particularly infancy and early adulthood) confirm that medical advancements have mostly reduced premature deaths, with less relative impact on the oldest age groups.
- **Mortality Trend (k_t):** This time-varying index tracks the overall direction of mortality. The steady downward trajectory confirms consistent historical improvements in life expectancy. The notable spikes (e.g., 1918, World War II) mirror the biometric shocks observed in our initial exploratory data analysis.




Figure 2: Baseline Profile (a_x)




Figure 3: Improvement Sensitivity (b_x)




Figure 4: Mortality Trend (k_t)

1.5 Old-Age Mortality: The Kannisto Method

A critical consideration in valuing pension liabilities is determining the ultimate age of the life table and deciding how to extrapolate mortality rates for extreme old ages. For this portfolio, the youngest participants are currently 30 years old. To ensure the complete liabilities over a 90-year projection horizon, the ultimate age is set to 120.

However, historical mortality data of the Netherlands at advanced age, for example 90+, is sparse for direct calibration. To illustrate this sparsity, we can observe the raw HMD death counts ($D_{x,t}$) from the year 2023. While middle-ages have relatively high sample sizes, data reduces exponentially at advanced ages:

```
nld_data %>% filter(Year==2023, Age %in% c(60,105,108))%>% select(Age,Total.x,Total.y)
Age      Total.x      Total.y
1 60      1396.00     249113.28
2 105       38.59        57.22
3 108        4.99         5.28
```

With only 9 person-years of exposure at age 108, calculating empirical central death rates becomes heavily distorted by statistical noise. To resolve this and smoothly complete the mortality table,

the Kannisto method is applied [3]. The Kannisto method assumes that at older ages, the logit transformation of the central death rate follows a linear relationship with age.

The procedure is implemented as follows:

1. **Estimation:** We isolate the reliably fitted Lee-Carter mortality rates for ages 80 to 90. An Ordinary Least Squares (OLS) regression is fitted to the logit of these rates:

$$\text{logit}(m_{x,t}) = \ln\left(\frac{m_{x,t}}{1 - m_{x,t}}\right) = \phi_1 + \phi_2 x \quad (3)$$

It is important to note that the logit function is mathematically defined only for inputs strictly between 0 and 1. If any raw central death rate equals or exceeds 1, the transformation fails. Because of this, we fit the regression using the stabilized rates from ages 80 to 90 rather than the volatile raw data at age 105+.

2. **Extrapolation:** Using the estimated coefficients ϕ_1 and ϕ_2 , the linear relationship is projected forward for ages 91 to 120.
3. **Transformation:** The inverse logit function is applied to the extrapolated values to return them to standard central death rates. Because the inverse logit function mathematically asymptotes at 1, it forces the extrapolated rates to be strictly bounded between 0 and 1. This provides a stable tail for liability valuation but it is a theoretical drawback of the Kannisto model: It artificially caps $m_{x,t}$ at 1. This prevents the rate from capturing real-world scenarios where mortality risk exceeds one death per person-year.

To validate this approach, the extrapolated Kannisto tail was plotted alongside the fitted Lee-Carter rates and the raw empirical HMD data (Ground Truth) for the year 2023. As visualized in Figure 5, the Kannisto extrapolation captures the trajectory of real-world mortality at advanced ages.




Figure 5: Mortality profile displaying the Lee-Carter fit (ages 80-90) and the Kannisto extrapolated tail (ages 90-120), plotted against the raw HMD ground truth data.

1.6 Evaluation of Period Life Expectancy

With a full matrix of $(m_{x,t})$ up to age 120, we can evaluate the baseline longevity risk of the pension participants using period life expectancy (e_x). In the calculation of period life expectancy, we assume frozen mortality profile. So, for example a 30-year-old today is assumed to experience the exact same mortality risk at age 80 as an 80-year-old experiences today.

The expected remaining lifetime is calculated by integrating the survival probabilities over the remaining lifespan. Assuming a piecewise constant force of mortality between integer ages, the period life expectancy is then computed from the projected central death rates.

Figure 6 illustrates the historical evolution of period life expectancy for several ages (35, 45, 55 and 65).




Figure 6: Historical evolution of period life expectancy for selected ages.

The visualization shows a steady, long-term upward trajectory in expected lifetimes across all observed ages. However, relying purely on period life expectancy for pension valuation structurally underestimates liabilities. Because, it fails to account for the future medical improvements that today's 30-year-old will experience as they age. This limitation necessitates the transition to cohort thinking in the subsequent phases of this analysis. Before moving to cohort based life expectancy in the next section, we conclude this section by calculating the expected present value (EPV) of a whole life annuity.

1.7 Financial Assumptions: EIOPA Risk-Free Rate

To accurately value the long-term pension liabilities, future cash flows must be discounted to their present value. For this purpose, we use the risk-free interest rate (RFR) term structures published by the European Insurance and Occupational Pensions Authority (EIOPA).

For this valuation, the baseline discount curve is derived from the EIOPA spot rates for the Netherlands (as of early 2026). Crucially, the curve incorporating the Volatility Adjustment (VA) is utilized. The reason why VA is utilized because pension liabilities are long-term. So, applying the VA prevents artificial fluctuations in the valuation caused by short-term volatility in bond spreads. The extracted term structure is visualized in Figure 7.




Figure 7: EIOPA Risk-Free Rate term structure for the Netherlands

1.8 Valuation of Period Liabilities

With both the closed mortality matrix and the EIOPA discount curve, we can calculate the Expected Present Value (EPV) of the pension liabilities.

The portfolio consists of lifelong pension benefits payable annually in advance, contingent upon the participant's survival. For a participant who has already reached the retirement age ($R = 65$), the EPV of a whole life annuity issued at age x in period t , with an annual payout C (e.g., 10,000 EUR), is mathematically defined as:

$$\ddot{a}_{x,t} = C \sum_{T=0}^{\infty} v(t, T)^T {}_T p_{x,t} \quad (4)$$

where:

- ${}_T p_{x,t} = \prod_{j=0}^{T-1} p_{x+j,t}$ is the T -year survival probability calculated under the period approach.
- $v(t, T) = \frac{1}{1+i(t, T)}$ is the annual discount factor for maturity T .

However, the ISSurance portfolio includes active participants as young as 30. For these individuals, the annuity is deferred. To accurately value this liability against the full EIOPA risk-free term structure, the valuation projects the potential cash flows across the participant's entire remaining lifespan, starting from their current age x .

An indicator function, $I(x + T \geq R = 65)$ is applied to ensure that pension payments only start once the retirement age is reached. Each individual future cash flow is then discounted back to the valuation date using the spot rate corresponding to its specific maturity T . The Expected Present Value (EPV) of the deferred annuity is mathematically defined as:

$$\text{Deferred EPV}_{x,t} = C \sum_{T=0}^{\infty} I(x + T \geq 65) \times v(t, T)^T \times {}_T p_{x,t} \quad (5)$$

This approach guarantees financial precision across the yield curve. For example, a cash flow paid to a current 30-year-old when they turn 70 (which occurs at $T = 40$) is discounted using the 40-year EIOPA spot rate, ensuring total alignment with institutional valuation principles.

Applying this logic across all portfolio ages where $C = 1$ yields the liability curve shown in Figure 8.




Figure 8: Expected Present Value (EPV) of a 1-unit deferred annuity across portfolio ages.

The curve demonstrates the non-linear relationship between age and expected liability. For younger participants, the EPV is heavily discounted due to the deferral period and the cumulative probability of dying before retirement. As participants approach age 65, the discounting effect diminishes. This drives the EPV to its peak. In the Post-retirement phase, the EPV declines as the expected remaining lifetime shortens. For a rigorous mathematical proof for this observation, see Appendix A

2 Period to Cohort Thinking

The period valuation in previous section relies on a static mortality profile. It assumes that future mortality rates will remain identical to the current year. However, the Lee-Carter trend parameter (k_t) proves that mortality consistently improves over time. To accurately reserve for future obligations, the valuation must be evaluated under cohort approach.

The difference between these two approaches is best visualized using a Lexis diagram (Figure 9).




Figure 9: Lexis diagrams illustrating period vs. cohort thinking [4].

The left panel represents period thinking. It takes a vertical slice of data. All mortality rates are extracted from a single, fixed calendar year. For example, it assumes a 50-year-old in 2026 will experience the exact same mortality risk at age 51 as an 51-year-old does today.

The right panel represents cohort thinking. It tracks a diagonal path. As a participant ages by one year, the calendar year also advances by one year. Under this framework, a 50-year-old in 2026 will experience the mortality risk of an 51-year-old in the year 2027. This correctly incorporates expected medical advancements of next year.

2.1 Cohort Life Expectancy

To implement this diagonal tracking, the `LE_period` function was modified into `LE_cohort`. While the period function was extracting mortality rates from a single row corresponding to the valuation year, cohort function extracts a diagonal vector from the matrix: $m_{x+k,t+k}$ for $k = 0, 1, 2, \dots$ until the ultimate age. Because this diagonal extends into the future, the function takes the projected Lee-Carter mortality matrix appended to the historical data as input.

This methodological shift alters the projected survival probabilities. Figure 10 illustrates this divergence by tracking the specific trajectory of a 30-year-old participant.

During the early decades of the projection (from age 30 to approximately age 60), the period and cohort curves remain nearly identical because baseline mortality risk during these working years is already low. However, as the participant approaches retirement and advanced old age, a substantial

gap emerges. Visually, the solid cohort curve is farther away from the dashed period curve and the steep drop in survival probability is shifted to the right. Practically, this means a 30-year-old has a higher likelihood of reaching ages 80 and beyond than a period life table would suggest. This increases the expected duration of pension payouts.




Figure 10: Survival probability curve (${}_tp_{30}$) for a 30-year-old participant.

Though, the impact of this shift depends heavily on age. Figure 11 compares period and cohort life expectancy across the portfolio. The region in between the cohort and the period curves represents the number of years resulting from using cohort approach instead of period approach. Younger participants gain several years of expected life. They have decades of future improvements ahead of them. Conversely, the gap for a 90-year-old is negligible. They have a very short horizon for future improvements to show their effect.




Figure 11: Period vs. Cohort Life Expectancy across portfolio ages.

To demonstrate that this divergence is a persistent feature, Figure 12 tracks the expected remaining lifetime of a 65-year-old (e_{65}) from 1970 to 2023. As visualized by the shaded region, the cohort life expectancy is always higher than the period life expectancy across the 50 years of observation window.

Furthermore, the period line exhibits fluctuations corresponding to specific annual shocks such as severe winter flu seasons or recent pandemics, whereas the cohort approach yields a much smoother

path as it integrates these shocks over the participant's remaining lifetime.




Figure 12: Historical evolution of life expectancy at age 65.

2.2 Valuation of Cohort Liabilities

The divergence between two approach impacts the financial valuation of the liabilities. To compute the cohort Expected Present Value (EPV), the underlying financial mechanics remain unchanged: an indicator function limits cash flows to ages 65 and above and the EIOPA term structure discounts each cash flow to the valuation date. The critical difference lies in the mortality input. The cohort valuation swaps static period survival probabilities for dynamic, diagonal survival probabilities (${}_t p_x^{\text{coh}}$). This isolates the precise financial impact of continuous mortality improvements.

Figure 13 maps the resulting EPV difference across ages for a unit annuity.




Figure 13: Period vs. Cohort EPV valuation across ages.

The shaded region shows the extra liability resulting from cohort approach. As the expected lifetime increases regardless of the participant age, the number of years that participants will be paid their annuity increases. However, this extra liability differs with the age. For participants with older ages like 80, 90, the two curves converge because their remaining lifespan is too short to experience meaningful medical advancements. To be mathematically rigorous, ${}_t p_x^{\text{coh}} \approx {}_t p_x^{\text{per}}$ when x is large. However, for younger participants (ages 30 to 50), the gap widens. Decades of compounding mortality improvements increase their probability of reaching and surviving in their retirement period.

Consequently, a valuation based on static period tables leaves the pension fund under-capitalized. To quantify the financial severity of this shortfall, Figure 14 translates the absolute EPV gap into a relative valuation difference. This metric shows the percentage by which the period method underestimates the true liability over time.




Figure 14: The historical under-reserving error

This visualization provides a risk management insight regarding age sensitivity. While the period method underestimates the liability for a 65-year-old retiree by a relatively small margin, the error compounds largely for younger active participants. For a 30-year-old in 1970, ignoring future mortality improvements resulted in mispricing the liability by around 25%.

Furthermore, tracking the metric historically reveals a downward trend across all age groups. As baseline mortality rates have improved from 1970 to 2023, the base probability of survival has increased. Consequently, the relative impact of future improvements on the total EPV has gradually decreased.

Despite this downward trend, the current shortfall remains relatively high. In 2023, period tables still underestimate the liability of a 30-year-old participant around 10%. This sensitivity shows the risk profile of ISSurance’s portfolios. Because the Green Portfolio is mostly concentrated in younger active members, relying on period tables could create a huge capital shortfall. Conversely, the older Gold Portfolio remains relatively riskless.

2.3 Longevity and Mortality Shocks and their Effects

To ensure adequate capitalization and deal with risks mentioned in the previous subsections, insurers must quantify their exposure to adverse demographic scenarios. For a pension portfolio, the primary risks are mortality and longevity shocks. While in the former, participants live longer than expected and this increases the total payout duration, in the latter the risk is participants die earlier than expected so this results in inefficient use of capital.

To evaluate this uncertainty, we apply two deterministic shocks to the baseline cohort mortality rates ($m_{x,t}$):

- **Longevity Shock:** A permanent 20% decrease in mortality rates ($m_{x,t} \times 0.80$).
- **Mortality Shock:** A permanent 15% increase in mortality rates ($m_{x,t} \times 1.15$).

We first evaluate the demographic impact of these shocks. Figure 15 illustrates the cohort life expectancy under the baseline, longevity and mortality scenarios.




Figure 15: Life Expectancy Risk across portfolio ages.

The longevity shock (red line) shifts the expected remaining years upward. This might represent a scenario of sudden medical breakthroughs. The mortality shock (green line) shifts the expectation downward. The shaded region between them defines the area of uncertainty.

We then translate this demographic uncertainty into financial liabilities. Figure 16 displays the Expected Present Value (EPV) of a unit deferred annuity under the same three scenarios.




Figure 16: Liability Risk Corridor showing the EPV under baseline and shocked scenarios.

For a pension provider, the longevity shock represents the adverse scenario. While the transition from period to cohort tables impacted younger participants more due to long horizon of projected medical advancements, a static 20% longevity shock impacts older participants more. Because baseline mortality is already extremely low for a 30-year-old, so a 20% reduction yields a negligible absolute increase in their probability of surviving to retirement.

However, for participants currently entering or within the payout phase (ages 65 to 75), the baseline mortality risk is much higher. This extends the payout duration precisely when the discount factor is highest. As a result, the absolute capital gap per unit of annuity peaks around age 70.

```

> df_mort_shocks %>%
  mutate(Abs_Diff = Long - Base) %>%
  arrange(desc(Abs_Diff)) %>%
  select(Age, Base, Long, Abs_Diff) %>%
  head(10)

```

	Age	Base	Long	Abs_Diff
1	70	13.00559	13.85579	0.8502058
2	69	13.47172	14.32079	0.8490715
3	71	12.53447	13.38340	0.8489241
4	68	13.93484	14.78183	0.8469923
5	72	12.06494	12.91192	0.8469767
6	67	14.39376	15.23723	0.8434783
7	73	11.59088	12.43392	0.8430365
8	66	14.84918	15.68800	0.8388159
9	74	11.11674	11.95395	0.8372143
10	65	15.29964	16.13241	0.8327728

Consequently, while the young Green Portfolio holds the highest longevity risk, the older Gold Portfolio holds the highest level longevity risk. So, the Gold Portfolio requires more capital buffers against deviations in mortality assumptions.

Interactive Mortality Shock Simulator

An interactive Mortality Shock Simulator was developed to visualize demographic risk assessment. Figure 17 displays the interface, which allows the user to dynamically adjust the mortality multiplier and observe the resulting impact on cohort life expectancy and expected present value across all ages.




Figure 17: Interactive Mortality Shock Simulator.

2.4 Financial Shocks: Interest Rate Sensitivity

Because pension liabilities are spread through several years, even decades, their present value is sensitive to discount rates. To assess this financial risk, we evaluate the liability across different economic environments using the EIOPA Risk-Free Rate (RFR) term structures released between 2016-2022 and 2026. Figure 18 illustrates the evolution of these yield curves.




Figure 18: EIOPA Risk-Free Rate Term Structure Evolution.

The macroeconomic environment shifted over the last decade. The years from 2017 to 2022 represent a period of financial suppression, with spot rates turning negative at the short end of the curve. Conversely, 2026 curve shows a higher-yield environment.

To quantify the impact of these shifts on the pension fund, Figure 19 evaluates the Expected Present Value (EPV) of the baseline cohort liability under each historical yield curve.




Figure 19: The impact of changing EIOPA discount curves on the EPV.

As expected, there is an inverse relationship between the interest rate and the liability valuation. The low-rate environments produce the highest liability valuations. This represents a considerable financial stress scenario for the company. On the other hand, the higher rates of 2026 heavily discount future cash flows, resulting in suppression in the present value of the obligations.

The vertical spread between the curves shows that the interest rate sensitivity is highly dependent on age. For older ages, the time of cash flows are not late, so the time-to-maturity (T) is short. This has minimal effect on the discount factor $(1 + i)^{-T}$, thus on the EPV. Consequently, EPV curves cluster on the right side of the chart. However, for a 30-year-old participant, the cash flows are deferred for a long time. So, the discount factor compounds the interest rate change over decades. Consequently, a drop in interest rates (e.g., shifting from the 2026 curve to the 2021 curve) inflates the present value of their deferred annuity, more than doubled in this case.

Lastly, the shape of the EPV curve depends on the yield curve. It depends on the rate level and convergence speed of curve to its ultimate rate level. The 2026 curve has high rate and also reaches the ultimate forward rate rapidly. This heavily discounts long-deferred cash flows. So, the result is a convex EPV profile. Conversely, the 2021 curve is suppressed, i.e., it takes longer time to converge. This diminishes the effect of discounting. In this case, valuation is driven primarily by cumulative survival probabilities (TP_x). This yields a concave (linear-like) EPV curve. Younger participants therefore become highly capital-intensive to reserve for.

This establishes a critical risk difference for ISSurance. While the older Gold Portfolio holds the highest level-risk against demographic longevity shocks, the younger Green Portfolio is more exposed to macroeconomic interest rate shocks.

2.5 Stochastic Interest Rates: Vasicek Model

The previous sections examined historical EIOPA data. We now simulate future interest rate paths to perform a financial stress test. The Vasicek model generates stochastic interest rate paths using a mean-reverting random walk. The Vasicek model, introduced in 1977 [5], generates stochastic interest rate paths by assuming rates follow a mean-reverting random walk. Four parameters determines the evolution of the curves: the initial short rate (r_0), the long-term mean level (θ), the speed of mean reversion (k), and the instantaneous market volatility (β).

Figure 20 displays five simulated yield curves.




Figure 20: Simulated spot rate curves using the Vasicek model ($n = 5, T = 100, r_0 = 0.025, \theta = 0.035, k = 0.15, \beta = 0.015$).

The random paths visualized here were generated earlier, unseeded global environment. Reproduction scripts are seeded to guarantee identical structural volatility, though simulated paths will vary.

We isolate Path 2 and Path 4 to highlight structural differences and evaluate their specific effects on liability valuation. Path 2 represents a short-term high-rate regime where the spot rate spikes above 6% near maturity year 10 before reverting downward toward the long-term mean. Path 4 represents a short-term suppressed regime where the spot rate drops below 1% in the first five years before increasing to stabilize above the mean for long maturities.

These distinct curve shapes impact the liability valuation, as illustrated by the EPV distribution across ages in Figure 21.




Figure 21: Cohort EPV evaluated under the simulated Vasicek interest rate scenarios.

The financial impact depends on cash flow timing. For a 60-year-old participant whose pension payments begin in five years, Path 4 provides extremely low rates during this near-term payout window. This lack of discounting inflates their EPV to the one of the highest level among all scenarios. Conversely, Path 2 features very high rates during this same window. So, the near-term cash flows become heavily discounted and this drives the EPV to the lowest level.

A 30-year-old participant faces a different financial reality because their cash flows are deferred for 35 years. Path 4 features high interest rates at these long maturities. This discounts the distant cash flows severely and generates the lowest EPV. On the other hand, Path 2 features lower rates at these same maturities, generating a comparatively higher EPV for the young participant.

These stochastic simulations reveal distinct portfolio vulnerabilities based on age composition. Older portfolios face severe valuation risks from short-term interest rate dips, making the Gold Portfolio highly exposed to regimes like Path 4. Younger portfolios face risks primarily from long-term interest rate declines, meaning the Green Portfolio requires hedging against these environments.

3 Stochastic Mortality Scenarios

3.1 Simulating Future Mortality

While deterministic shocks provide a standard baseline for stress testing, they lack statistical grounding. They cannot assign explicit probability metrics to extreme demographic events. To quantify true tail risks, we apply stochastic modeling. We simulate 1,000 future paths for the Lee-Carter mortality trend parameter (k_t).

This simulation relies on the historical time-series properties of k_t . The model projects future values by assuming a random walk with drift. It generates distinct paths by adding normally distributed random errors at each annual time step.

Figure 22 displays this stochastic evolution. As the projection extends further into the future, the simulated paths fan out. This visually represents the compounding uncertainty of mortality improvements over time. Each individual path constitutes a future mortality scenario. To ensure every simulated life table spans the entire portfolio horizon, we apply the Kannisto method to close the rates up to age 120.




Figure 22: Stochastic evolution of the mortality trend (k_t) across 1,000 simulations.

3.2 Life Expectancy Distributions

Using these simulated mortality surfaces, we calculate the cohort life expectancy under each scenario for all portfolio ages. Figure 23 illustrates the probability density of these expectations.

The distribution shapes vary significantly by age. Younger cohorts exhibit much wider distributions. Their long projection horizons allow the stochastic variance in the k_t parameter to compound over years. Conversely, older cohorts have a shorter remaining lifespan. This limits the impact of future mortality uncertainty. This results in narrow, highly concentrated density peaks.




Figure 23: Life Expectancy distributions across portfolio ages.
An interactive, 3D web version of this plot can be found [here](#).

To compare stress scenarios to our stochastic simulations, we isolate the 50-year-old cohort. Figure 24 plots the simulated life expectancy distribution against the deterministic shocks calculated previously. The distribution roughly follows a normal bell curve centered around the base cohort estimate.




Figure 24: Distribution of Life Expectancy for a 50-year-old against deterministic shocks.

We define the 2% and 98% Value at Risk (VaR) boundaries to represent severe demographic events. The yellow vertical lines indicate these stochastic limits. The red and green lines represent the static Solvency II shocks.

We observe that the 98% VaR aligns almost perfectly with the 0.8x longevity shock. This establishes a clear relationship for the right tail of the distribution. The deterministic multiplier acts as an accurate proxy for extreme longevity risk.

However, the opposite tail reveals a distinct discrepancy. The 1.15x mortality shock falls short of the 2% VaR boundary. It is even closer to the base mean. It fails to capture the full extent of the stochastic downside risk. This indicates an asymmetric relationship between the two methodologies. The static framework captures extreme life extensions for this cohort. Conversely, it underestimates the severity of early mortality events modeled by simulations.

3.3 Stochastic Liability Valuation

To translate these demographic distributions into financial liabilities, we calculate the Expected Present Value (EPV) of a deferred annuity for every simulated path. We utilize the base 2026 EIOPA yield curve. Figure 25 maps the resulting EPV distributions across all ages.

The financial risk profiles are similar to the demographic risk profiles. Younger cohorts exhibit much wider EPV distributions. Their long deferral periods compound the stochastic mortality variance. Older cohorts exhibit narrow, highly concentrated liability peaks.




Figure 25: Expected Present Value distributions across portfolio ages.
Note: An interactive, 3D web version of this distribution can be found [here](#).

Figure 26 isolates the financial distribution for the 50-year-old cohort. We compare the stochastic VaR boundaries against the financial liabilities calculated under the deterministic shocks.




Figure 26: Distribution fitting of Expected Present Value for a 50-year-old.

The right tail of this distribution represents severe longevity risk. Extended lifespans directly inflate

the present value of the annuity. The 98% financial VaR aligns well with the liability calculated under the 0.8x longevity shock. The standard deterministic multiplier successfully prices the extreme longevity risk for 50 year-old.

The left tail represents early mortality where the expected payout duration decreases. The liability under the 1.15x mortality shock fails to reach the 2% VaR boundary, so it underestimates the financial impact of extreme downside mortality scenarios. This is in line with the identical asymmetric relationship observed in the life expectancy analysis.

3.4 Age-Dependent Capital Adequacy

While deterministic shocks provide an approximation for specific cohorts like the the 50-year-old above, their validity across the demographic spectrum differ. Figure 27 visualizes the percentage gap between the static multipliers and the internal stochastic VaR boundaries across all portfolio ages.




Figure 27: Age-Dependent Adequacy of Regulatory Shocks compared to Stochastic VaR.

The empirical data shows an age-dependent drift. Focus first on the longevity risk (the red curve). For a 30-year-old, the 0.8x shock dangerously under-reserves capital, falling roughly 4.5% short of the 98% stochastic VaR. Young cohorts have decades for medical improvements to compound. This surpasses the linear impact of a static multiplier. Conversely, for a 90-year-old, the same shock over-reserves the liability by roughly 8.0%. It assumes an unrealistic acceleration of extreme old-age survival resulting in requirement of excess capital.

On the other hand, the green curve exhibits the inverse relation. It highly overestimates downside volatility for younger cohorts, sitting more than 10% above the true stochastic boundary. However, it fails to reach the 2% VaR boundary for participants over age 75. This underestimates the financial impact of mortality shocks in old age.

Importantly, the two curves intersect the zero-gap line at different ages. There is no single representative cohort where both deterministic shocks provide accurate capital metrics simultaneously.

This proves that a single static multiplier is not a necessary tool. It does not necessarily reflect the true risk profile of a pension fund. The standard shocks under-price risk for young cohorts and over-price it for older ones. Because of this, the total capital requirement depends entirely on the fund's specific age distribution. So, we cannot measure solvency by looking at isolated age groups. We must combine the cash flows of the entire population using our simulated mortality paths.

4 Portfolio Valuation

4.1 Individual Cashflow Projections

To value the aggregate liabilities, we first project the expected cashflow vector for an individual participant. We define the expected payout at any future time T as the base annuity amount multiplied by the probability of surviving to that time, activated when retiree reaches the age of 65.

Figure 28 visualizes this cashflow evolution, an active 50-year-old participant against a retired 70-year-old participant.




Figure 28: Expected cashflow projections deferred (Age 50) and immediate (Age 70) annuities under 100 stochastic mortality paths.

The structure of the pension liability is set by the participant's current age:

The Active Cohort (Age 50): This represents a deferred annuity. The expected payout remains zero during the 15-year deferral period. At age 65, the expected cashflow spikes, reflecting the cumulative probability of surviving to retirement. The surrounding stochastic fan illustrates demographic uncertainty. Because a 50-year-old has decades for medical improvements to compound before reaching old age. That's why the dispersion in the tail is wide.

The Retired Cohort (Age 70): This represents an immediate annuity. The participant is already over age 65. This means the cashflow initiates immediately in the valuation year at exactly 1.00 EUR (100% survival probability at $T = 0$). From there, the expected payout immediately begins a steady, continuous decay as natural mortality reduces the cohort.

By applying this cashflow mechanism to every age group and multiplying by the respective participant counts, we map the liquidity requirements for the portfolios over the next years.

4.2 Aggregate Portfolio Cashflows

By evaluating the individual cashflow vectors alongside the demographic composition of each portfolio, we can project the total aggregate liquidity requirements. Figure 29 illustrates the best estimate of these aggregate cashflows over the next 100 years.




Figure 29: Aggregate expected cashflows for the Green, Silver and Gold portfolios.

The sawtooth-like pattern occurs because the portfolio data is grouped into 10-year age cohorts. Each sudden spike represents a younger age cohort reaching the retirement age of 65 and initiating their annuities. For instance, the spike at year 2031 corresponds to the payment of the people who are currently at age 60.

More importantly, the graph shows different risk profiles across the portfolios:

- **The Gold Portfolio (Liquidity Risk):** As gold portfolio represents an aging demographic, its peak cash requirements occur almost immediately. While the total liability runs off rapidly, it requires highly liquid, short-term assets to fund the immediate outflows.
- **The Green Portfolio (Duration Risk):** As green portfolio is dominated by younger active members, its cashflow requirements are heavily deferred. While it demands less immediate liquidity compared to others, its liability tail extends most into the future. This creates duration risk and it requires long-term asset matching to protect against fluctuating interest rates over the next 80 years.

4.3 Total Liability Valuation and Uncertainty

To calculate the Present Value (PV) of these liabilities, we discount the aggregate cashflows using the base 2026 EIOPA risk-free rate curve. However, a single deterministic best estimate is insufficient for capital reserving.

To quantify demographic uncertainty, we evaluate the total portfolio liabilities across all 1,000 stochastic mortality paths generated in Phase 3. Figure 30 displays the probability distributions of the total present values.

This visualization provides the following insight into demographic risk profiling: The portfolios do not just differ in total cost but also differ in uncertainty.

- **The Gold Portfolio (High Predictability):** Exhibits a narrow and highly concentrated distribution. Because its participants are already at advanced ages, the portfolio is in a fast runoff phase. The short remaining time horizon prevents future mortality trends from compounding. This results in a highly predictable liability with minimal tail risk.
- **The Silver Portfolio (Transitional Risk):** Serves as the demographic bridge. Its distribution is wider than Gold's but more contained than Green's. It faces a blended risk profile, balancing liquidity needs with moderate exposure to long-term longevity improvements.

- **The Green Portfolio (Severe Tail Risk):** Exhibits a highly dispersed, right-skewed distribution. Young participants have decades of life expectancy remaining. Any systemic improvement in mortality compounds a lot over their lifetimes. This risk pushes the right tail of the Green distribution severely outward.




Figure 30: Distribution of Total Present Value for each portfolio across mortality scenarios.

4.3.1 Capital Requirements and Risk Margin

Consequently, while the Green portfolio holds the largest absolute liability, it also carries the highest relative variance. ISSurance cannot apply a uniform capital charge across the company. The distance between the deterministic Best Estimate and the 98% Value-at-Risk (VaR) defines the risk margin.

- **Green Portfolio:** Best Estimate of 73,189 EUR; 98% VaR of 77,618 EUR (Ratio: **1.061**).
- **Silver Portfolio:** Best Estimate of 68,778 EUR; 98% VaR of 72,589 EUR (Ratio: **1.055**).
- **Gold Portfolio:** Best Estimate of 64,368 EUR; 98% VaR of 67,605 EUR (Ratio: **1.050**).

The VaR-to-Best-Estimate ratio serves as an approximation for the capital intensity of each portfolio. The Green portfolio requires a 6.1% risk margin to cover 1-in-50-year longevity events, whereas the Gold portfolio requires only 5.0%. This 1.1% gap in relative risk proves that the tail of a younger demographic has a measurable impact on capital efficiency.

4.4 Sensitivity Analysis: Biometric vs. Financial Shocks

To fully assess portfolio risk, we evaluate how the Total Present Value reacts to standard regulatory and financial shocks. We focus on the percentage change relative to the base cohort valuation to properly compare the portfolios.

4.4.1 Longevity and Mortality Shock

We first apply the deterministic multipliers: a 0.8x factor for a longevity shock and a 1.15x factor for a mortality shock. Figure 31 illustrates the relative impact of these scenarios.

The relative impact of the biometric shocks is almost uniform. The longevity shock increases the liability by 6.5% to 6.7% across all portfolios, while the mortality shock decreases the liability by an identical 4.2%.




Figure 31: Relative impact of Solvency II biometric shocks on Total Present Value.

Given that younger cohorts have decades of future projection time, one might expect the Green portfolio to show a much higher percentage sensitivity to longevity. However, this biological exposure is neutralized by the time value of money. For the Gold portfolio, extended life expectancies generate extra cash flows in the near term. These face minimal interest rate discounting. For the Green portfolio, the extra survival years occur 50 to 60 years in the future. The heavy discounting applied to these distant cash flows cancels out the compounding mortality effect.

4.4.2 Interest Rate Sensitivity

Next, we evaluate financial sensitivity by applying parallel shifts of +100 and -100 basis points to the base EIOPA risk-free rate curve. Figure 32 displays the results.




Figure 32: Relative impact of parallel interest rate shifts on Total Present Value.

Unlike the mortality shocks, the financial stress exhibits huge differences between the portfolios. The Green portfolio is highly sensitive to interest rate movements, with its liability swinging between +20.5% and -15.7%. In contrast, the Gold portfolio shows lower sensitivity (+15.0% to -12.0%).

This variance is driven by duration risk. An interest rate shock directly alters the discount factor $v^T = (1 + i)^{-T}$. Because the Green portfolio consists of deferred annuities, its cash flows occur much later (larger T). A 1% shift in the interest rate compounds exponentially over these longer time horizons. In conclusion, while the longevity shocks impact all portfolios more or less equally in relative terms, the Green portfolio is exposed to financial market volatility the highest.

5 Conclusion and Recommendations

This report evaluates the liabilities of the company across different portfolios. By transitioning from static period life tables to dynamic cohort projections and applying stochastic simulations, risk factors are identified. To ensure solvency, the following actions are recommended:

1. Cohort-Based Reserving

Relying on period life expectancy structurally underestimates future liabilities. Because the Green portfolio consists primarily of young active members, ignoring future mortality improvements results in a capital shortfall. ISSurance must adopt cohort-based valuation for reserving to account for medical advancements.

2. Duration Risk in the Green Portfolio

The Green portfolio represents the largest long-term liability. Sensitivity analysis proves that its deferred cash flows make it highly vulnerable to interest rate fluctuations. A 100 bps drop in interest rates inflates its liability by over 20%. ISSurance must implement an asset matching strategy like investing in long-term bonds to hedge against if yield curves fall.

3. Liquidity Risk in the Gold Portfolio

The Gold portfolio represents an aging demographic. While it is relatively protected from interest rate shifts, it requires immediate cash outflows. The investment strategy for this portfolio must prioritize liquid, short-term assets to fund continuous annuity payments without forcing the sale of assets.

4. Dynamic Risk Margins

The stochastic simulations revealed that uncertainty is not uniformly distributed. The Green portfolio exhibits a wider dispersion of tail risk compared to the Gold portfolio. ISSurance should move away from flat capital charges and apply age-dependent risk margins that reflects the compounding mortality uncertainty of younger demographics.

By aligning the asset allocation strategy with the distinct duration and liquidity profiles of these three portfolios, ISSurance can secure its obligations against both demographic and financial shocks.

A Proof of Peak Expected Present Value at Retirement Age

For a participant aged $x < 65$, the expected present value of the deferred annuity is defined as $EPV(x) = v^{65-x} \cdot {}_{65-x}p_x \cdot \ddot{a}_{65}$. The relationship between consecutive ages can be evaluated using the backward recursion formula: $EPV(x) = v \cdot p_x \cdot EPV(x+1)$. Assuming a strictly positive interest rate ($i > 0$), which is the case for the Netherlands, the one-year financial discount factor $v = (1+i)^{-1}$ is strictly less than 1. Because the one-year survival probability $p_x \leq 1$, their product must satisfy $v \cdot p_x < 1$. Substituting this into the recursion yields $EPV(x) < EPV(x+1)$, demonstrating that the liability monotonically increases each year during the deferral period.

Conversely, for a participant in the payout phase ($x \geq 65$), the liability simplifies to an immediate whole life annuity: $EPV(x) = \ddot{a}_x = \sum_{t=0}^{\infty} v^t \cdot {}_t p_x$. Under standard actuarial mortality dynamics, the force of mortality strictly increases with age. This implies that for any duration $t > 0$, the cumulative survival probability of a younger individual is greater than that of an older individual, meaning ${}_t p_x > {}_t p_{x+1}$. Because the financial discount curve v^t applies identically across both summations, the term-by-term comparison dictates that $\sum v^t {}_t p_x > \sum v^t {}_t p_{x+1}$, yielding $EPV(x) > EPV(x+1)$. Because the sequence monotonically increases for $x < 65$ and monotonically decreases for $x \geq 65$, the global maximum of the expected liability mathematically occurs exactly at $x = 65$, retirement.

B Source Code and Reproducibility

For reproducibility, the R codebase is available in the link below.

Github Repository: github.com/keremiis/ALIM-HW

- **Phase 1 (Historical Data & Period Valuation):** HMD data, Lee-Carter model calibration, Kannisto tail extrapolations and period life expectancy calculations.
[\[phase_1.R\]](#)
- **Phase 2 (Cohort & Deterministic Shocks):** Transition to cohort survival probabilities, historical interest rate shocks and Vasicek interest rate scenarios.
[\[phase_2.R\]](#)
- **Phase 3 (Monte Carlo Simulations):** Monte Carlo simulations for the Lee-Carter k_t parameter, 3D density distributions and VaR checks.
[\[phase_3.R\]](#)
- **Phase 4 (Portfolio Valuation):** Expected cashflow generation, aggregate portfolio liabilities, 98% VaR metrics and sensitivity analysis.
[\[phase_4.R\]](#)
- **Interactive Dashboard:** Application for mortality shock simulation.
[\[app.R\]](#)

References

- [1] Lee, R. D., and Carter, L. R. (1992). Modeling and Forecasting US Mortality. *Journal of the American Statistical Association*, 87(419), 659–671.
- [2] Human Mortality Database. Available at <https://www.mortality.org/Country/Country?cntr=NLD> (data downloaded on April 9, 2026).
- [3] Kannisto, V. (1992). *Presentation at a workshop on old age mortality*. Odense University, June 1992.
- [4] Antonio, K., Devriendt, S., et al. (2017). Producing the Dutch and Belgian mortality projections: a stochastic multi-population standard. *European Actuarial Journal*, 7(2), 297–336.
- [5] Vasicek, O. (1977). An equilibrium characterization of the term structure. *Journal of Financial Economics*, 5(2), 177–188.